

The intellect of man is forced to choose

# Python for Poets

~~YOUR CODE IS BAD AND YOU SHOULD FEEL BAD~~

Your algorithm choice could be better

<https://medium.com/capital-one-tech/heat-death-of-the-universe-and-faster-algorithms-using-python-dict-and-set-f31517e7fa76>



# Topics

- Floating point and integer numbers and their ranges/precisions
- Gross Orders of Complexity
- Important Rules
- Efficient structures for searching, counting, analytics
  - Building fast indices for searching
- Case Study: simple analytics using Python
  - AKA: "The Poetry Corpus"



Background — On Computing



# Things that Matter

- ◆ Python is written in C
- ◆ NumPy is written in C (mostly)
  - ◆ Uses bare-metal data types
  - ◆ Imposes a number of limitations
- ◆ Pandas uses NumPy



# Things that don't matter

## Automatic Optimization

- ◆ Python doesn't optimize
- ◆ If it did debugging would be impossible

## Tricks

- ◆ There are no secret "turbo-boost" coding tricks
- ◆ Use packages like numpy, pandas or numba



# Integers

- ◆ Python: no limits — two internal representations
  - ◆ (fast) 64-bit values
  - ◆ (slow) Arrays of values (base  $2^{30}$ ) for numbers  $> 2^{64}$
- ◆ Numpy: int8, int16, int32, int64



# Float

$$f \approx [m_{10}] \times 10^{e_{10}} \approx [m_2] \times 2^{e_2}$$

- ◆ Python: IEEE 754 compliant 64-bit
  - ◆ 64-bit values. Source text is decimal 6.0221409e+23
  - ◆ Internal representation is binary  $\frac{1,121,711,153,537,035}{2^{50}} \times 2^{79}$
  - ◆ Mantissa is ~48 bits = ~16 digits. Exponent is  $\pm 300$
- ◆ Numpy: float32, float64



# Float is an Approximation

Infinite Repeating "binary point" Values...  
Are truncated

>>> 100+1/3-100

0.3333333333333333286

>>> 100+1/3-100-1/3

-4.718447854656915e-15

Approximation doesn't match Abstraction



This is going to involve work, right?



# Orders of Complexity

- $O(1)$  ♦ constant time
- $O(n)$  ♦ scales linearly with the amount of data
- $O(\log_2 n)$  ♦ scales with the log of the data. This is almost always because of some clever divide-and-conquer search
- $O(n \log_2 n)$  ♦ sorting and similar algorithms that do repeated searches
- $O(n^2)$  ♦ compare every item with every other item
- $O(2^n)$  ♦ whoa! The powerset of all subsets
- $O(n!)$  ♦ combinatoric explosion — all combinations of items

These are bad





# Worst Case — Permutations

$O(n!)$  Permutations — All Possible Orderings

For  $n=10$  items, there are 3,628,800 orderings

This is why we have sophisticated approximation-based algorithms

Optimal solutions would take centuries to find



# Really Bad Case — Power Set

$O(2^n)$  powerset — set of all subsets

For a  $n=10$  data set, there are 1,024 different subsets

Let's say you it takes 3 seconds to fetch one of the subsets

Overall? 51 minutes.



# A Bad Case — Comparisons

$O(n^2)$  matrix — compare each item against every other item

For  $n=10,000,000$  row dataset, that's  $10^{12}$  operations

Let's say comparison each takes a looong 1 ms ( $10^{-3}$  sec)

So.  $10^9$  seconds

= 32 years

On a 64-core processor, it's only 6 months!



# Not Too Bad — Sorting

$O(n \log_2 n)$  sort — each item does a  $\log_2 n$  lookup

For a  $n=10,000,000$  row dataset, that's  $2.3 \times 10^8$  operations

In memory, and each operation takes 0.01 ms ( $10^{-5}$  sec)

= 40 minutes



# Important Rules



First Rule of Optimization:

Don't



“We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%.”

—Donald Knuth



# Don't optimize something that doesn't work

- ◆ You have an app that's slow
- ◆ You must have rock-solid unit test cases
- ◆ Until you have rock-solid unit test cases, DO NOT OPTIMIZE

"it's no sin for an optimizing compiler to make a wrong program worse" — Bill McKeenan



# Don't Optimize Until You Profile

- ◆ The Pareto Principle
- ◆ Most of your program (80%) is fine
- ◆ Some small part (20%) of your program uses most of the resources
  - ◆ Hint: It's inside a loop somewhere
- ◆ Find the one function that's doing the most work; fix only that




# Tools Summary

AFTER you get the algorithm right



# Some Tools

- ♦ `pytest` — You must have unit test cases. You must have code coverage on the things you're going to optimize
- ♦ `logging` — Generally identify the likely location of problems
- ♦ `profile` — Pinpoint a “hot spot” where performance is really bad
-  ♦ `timeit` — Explore alternatives to find one that's fastest
- ♦ `sys.getsizeof()` — Some sense of the size of an object. See <https://docs.python.org/3/library/sys.html#sys.getsizeof>
- ♦ `%prun` — IPython profiler (there are others)
- ♦ `%time` — IPython timing
- ♦ `%%timeit` — rich timing details



# Algorithm and Data Structure

Or

How Do I Avoid  $O(n^2)$  ?



# Two species of algorithms

Search & Sort —  $\exists$  — There Exists

- ◆ This is where we often wind up with  $O(n^2)$  (or worse) kinds of problems
- ◆ Algorithm and data structure matters

Everything Else —  $\forall$  — For All (map and reduce)

- ◆ This is mostly bulk data movement — \*should\* be  $O(n)$
- ◆ Memory matters



# General Approaches

- ◆ Sets.  $O(1)$  Lookup. Size of set doesn't matter
- ◆ Dictionaries.  $O(1)$  Lookup. Size of dict doesn't matter
- ◆ The `bisect` module had  $O(\log_2 n)$  search of a `sorted()` list



# Some Examples

- ◆ Impossible  $O(2^n)$  problem
- ◆ Case Study: Joins (“lookups”) between two structures
  - ◆ We’ll look at stop-words lookup
  - ◆ We’ll look at bag-of-words vectorizing



# Impossible Problem — Profile All You Want

$$\sum_{x:x \in S} = t$$

Search a set,  $S$ , for a subset,  $s$ , with a given total,  $t$ .

given  $S = \{a, b, c, \dots\}$

find  $s \subset S$  where  $\sum s = t$



# Simple, Elegant, Unscalable

```
from itertools import chain, combinations
```

```
def powerset(iterable):
```

```
    "powerset([1,2,3]) -> () (1,) (2,) (3,) (1,2) (1,3) (2,3) (1,2,3)"
```

```
    s = list(iterable)
```

```
    return chain.from_iterable(combinations(s, r) for r in range(len(s)+1))
```

```
def exact_sum(s, target):
```

```
    for subset in powerset(s):
```

```
        if sum(subset) == target:
```

```
            return subset
```

$O(2^n)$



# Case Study



- ◆ Cleaning and Vectorizing Text
- ◆ Goal is a word vector for the “interesting” words
- ◆ Filter out stop words like “the” and “and”
- ◆ This reflects a common design pattern where there’s a lookup from one data structure to items in another
  - ◆  $O(n \times m)$  if we’re not careful
- ◆ There’s also a parsing aspect to decompose poetry lines to words.



# The Poetry Data

<https://github.com/aparrish/gutenberg-poetry-corpus>

- ◆ <http://static.decontextualize.com/gutenberg-poetry-v001.ndjson.gz>
- ◆ [https://raw.githubusercontent.com/nltk/nltk\\_data/gh-pages/packages/corpora/stopwords.zip](https://raw.githubusercontent.com/nltk/nltk_data/gh-pages/packages/corpora/stopwords.zip)



# A Newline Delimited JSON Reader

```
poetry_path = Path.cwd()/"guttenberg-poetry-v001.ndjson.gz"
def poetry_line_iter(source_path: Path=poetry_path) -> Iterator[str]:
    """Read the GZIP file via a decompressor."""
    with gzip.open(source_path) as source:
        for line in source.readlines():
            json_line = json.loads(line)
            yield json_line['s']
```



# Survey the File

```
for p in poetry_line_iter():  
    print(p)
```



# Reading the GZIP File

- ◆ Less I/O (fewer physical pages of data)
- ◆ More computation
- ◆ Is it worth it?



# A Stopword Iterator

```
sw_path = Path.cwd()/"stopwords.zip"

def stopword_iter(source_path: Path=sw_path) -> Iterator[str]:
    """Read a tiny subset of the ZIP file."""
    with zipfile.ZipFile(source_path) as archive:
        with archive.open("stopwords/english") as words:
            for line in words:
                yield line.decode('ascii').rstrip()

    yield from (
        "thy", "thou", "thee", "thus", "oh", "hath", "tis", "us", "forth",
        "thus", "ye", "shall", "thine")
```



# Survey the File

```
for s in stopword_iter():  
    print(s)
```



# Normalizing Words

- ◆ We'll strip almost all punctuation
- ◆ We can't strip all punctuation — we're vs. were
- ◆ Hyphenated words have single-word semantics
- ◆ Multiple apostrophes are rare (fo'c'sle, for example)



# Words from Each Line

```
def word_iter(text: str) -> Iterator[str]:  
    words = re.compile(  
        r"[a-z]+['''][a-z]+|[a-z]+(?:-[a-z]+)+|[a-z]+" )  
    for m in words.finditer(text.lower()):  
        yield m.group(0)
```



# The Regular Expression

- ◆ Kind of complicated
- ◆ A lot of computation
- ◆ Is it worth it?
  - ◆ It turns out, it's hard to do better than this
  - ◆ Feel free to try



# Survey The File (again)

```
for line in (  
    list(word_iter(text))  
    for text in poetry_line_iter()  
):  
    print(line)
```

Second

First



- ♦ ['to', 'lonely', 'hamlet', 'and', 'to', 'stirring', 'town']
- ♦ ['cheering', 'the', 'wayworn', 'traveller', 'as', 'it', 'flows']
- ♦ ['when', 'all', 'the', 'fields', 'with', 'drought', 'are', 'parched', 'and', 'bare']



# Remember The Intro?

- ◆ The biggest problem is Search
- ◆ Searching for Stopwords
- ◆ Searching for Vectorization words
- ◆ Search, Search, Search

Then nowise worship dusty deeds,  
Nor seek, for this is also sooth,

Look for  $O(n)$   
Try to replace with  $O(\log n)$  or  $O(1)$



# Removing Stopwords

```
def stopword_filter_1(stopwords: Iterable[str],  
words: Iterable[str]) -> Iterable[str]:
```

```
    sw_list = list(stopwords)
```

```
    for w in words:
```

```
        stop = False
```

```
        for sw in sw_list:
```

```
            if w == sw:
```

```
                stop = True
```

```
                break
```

```
    if not stop:
```

```
        yield w
```

Explicitly

$O(n \text{ poetry words} \times m \text{ stop words})$

Can we do better?



# Stopwords 2

```
def stopword_filter_2(stopwords: Iterable[str],  
words: Iterable[str]) -> Iterable[str]:
```

```
    sw_list = list(stopwords)
```

```
    for w in words:
```

```
        if w not in sw_list:
```

```
            yield w
```

What's  $O(x)$  of 'in'?



# Stopwords 3

```
def stopword_filter_2(stopwords: Iterable[str],  
words: Iterable[str]) -> Iterable[str]:
```

```
    sw_list = set(stopwords)
```

```
    for w in words:
```

```
        if w not in sw_list:
```

```
            yield w
```


What's  $O(x)$  of 'in'?



# set v. list

- ♦ `timeit.timeit("'a' in stop",  
 setup="from string import printable;  
 stop=set(printable)")`


- ♦ `Out[28]: 0.028795376998459687`



$O(1)?$

- ♦ `timeit.timeit("'a' in stop",  
 setup="from string import printable;  
 stop=list(printable)")`

- ♦ `Out[29]: 0.1707196079987625`



$O(n)?$



# Hidden Alternative #4

- ◆ The `bisect` module

```
array = sorted(items)
```

```
array[bisect.bisect(array, value) - 1] == value
```



# How to Avoid Search

- Set does hash-based lookup  $O(1)$
- Dict also hash-based  $O(1)$
- bisect is tree-like  $O(\log n)$





Next Steps



# Combining Things

```
word_bag = Counter(  
    stopwords_filter_1(  
        stopwords_iter(), First  
        poetry_word_iter(poetry_line_iter())  
        third second  
    )  
)
```



# Survey

```
word_bag = Counter(  
    stopwords_filter_2(  
        stopwords_iter(),  
        poetry_word_iter(poetry_line_iter())  
    )  
)  
  
word_vector = sorted(  
    k for k, v in word_bag.most_common(128))  
  
print(word_vector)
```



# Vectorize

```
def vectorize_3(word_vector: Sequence[str], line: str)
-> List[int]:

    index_iter = (
        (bisect(word_vector, w)-1, w)
        for w in word_iter(line))

    valid_index_set = set(i
        for i, w in index_iter if word_vector[i] == w)

    return [1 if i in valid_index_set else 0
            for i in range(len(word_vector))
            ]
```



# Conclusion

- ◆ Know the access cost for your data structure
- ◆ Get to  $O(1)$  or  $O(n)$  whenever possible
- ◆ Avoid  $O(n \times m)$
- ◆ Compress Data, use Generator Expressions



# Appendix



# The Big Picture — Memory

less bulk means more cache



# Tiers of Memory

| Memory        | Size            | Speed                 | Applications                                 |
|---------------|-----------------|-----------------------|--|
| Network/Cloud | Vast            | Glacial (s)           |  |
| Database      | Large (Tb - Pb) | Very Slow (ms)        |  |
| Local Disk    | Large (Tb - Pb) | Slow ( $\mu$ s to ms) |  |
| RAM           | Small (Gb - Tb) | Fast (ns to $\mu$ s)  | Large data structures                        |
| Cache         | Tiny (Kb - Mb)  | Superfast (ns)        | Small data structures, numbers, small tuples |



# Not All RAM Is The Same

- ◆ Hardware RAM is the actual memory transistors
- ◆ Swap Space is RAM-like storage on a local disk drive
- ◆ On small systems (under 32Gb) there will often be swap
  - ◆ OS will shuffle pages in and out of physical RAM
  - ◆ Dirty pages must get re-written and are expensive
  - ◆ Code pages are read-only and are cheap
- ◆ Over 32 Gb of RAM? swap has few benefits



# Consequence 1 — Compression

How much data are we talking about? Mb, Gb, Tb?

Where is the data? cache, RAM, file, database?

Compressed data (avro, parquet, etc.)

- ◆ The time decompress is (often) less than the time to transfer
- ◆ The cost to compress can be high
- ◆ You amortize a high cost to write against low cost to read **many** times



# Consequence 2 — Cache-only Data

- ◆ A “lazy” processing pipeline
- ◆ Python’s `map()` and `filter()`
- ◆ Generator expressions fit in cache naturally
  - ◆ Does not bring large volumes of data into RAM
  - ◆ May run mostly in cache
- ◆ You can’t force this... But you can encourage it by keeping less in memory



# Why this is hard

A common kind programming example reduces to

```
df = pandas.read-something()
```

struggle with df

This may not always work out well

The data is in RAM — which is better than disk or network — but isn't as good as cache

Reducing the volume of data is essential